

FIELD STUDY PREPARED FOR NORTH COUNTRY CHARTERS

Your FAQ asks guests to phone because voicemails and emails aren't always seen, and real money rides on that.

HOMER, ALASKA 2 AUGUST 2026

IN SHORT

Your FAQ carries one paragraph that does two things at once. It opens with "We have phenomenal office staff and we do everything we can to resell cancelled seats," and it ends with "IF YOU NEED TO CANCEL please call us and confirm over the phone with us that you are cancelling. We don't always see voicemails or emails right away and every minute counts."

Both halves are true together, and the second is not a complaint about the first. Your office can be as good as you say and a message can still sit unseen, because those are different problems. One is about people, the other is about a channel, and this is about the channel. Nothing you publish says the phones ring off the hook, so we won't imply it.

We read everything you publish, then five other Homer operators, Central Charters, Smokey Bay Air, Emerald Air Service, Destination Alaska Charters and Bear Viewing Alaska. On booking you are level with them, not behind, and none of them runs AI guest response either.

What we'd build first has no AI in it. A plain automatic reply on your shared inbox carrying your own FAQ sentence, so every guest who writes in gets your call-and-confirm request instead of only the ones who read the FAQ page. A Gmail setting, ten minutes of work.

What we'd ask you to buy is two weeks of counting. Two thousand eight hundred dollars fixed, plus about eight hours of your time. Nothing logs how many messages reach you or how many are about cancelling, so no honest price exists for the real build. The fortnight produces that number and ends with a verdict that includes "build nothing further."

On the money, the conservative case gets back about what the fortnight costs, the most likely case about twice it. Two conditions before you believe either. It only counts if you'd otherwise have bought the bigger build blind, and credit only the first year of that build and the conservative case recovers 83 percent, which doesn't clear.

We couldn't verify your charter rates or your message volume, so the bigger build gets a break-even you can test, not a return we invented.

The ask is a reply. Not a call.

01

You already wrote down where this hurts, and you wrote the answer to it too

The strongest thing we found about your business, you published yourselves.

One number. One shared inbox. Hours of "Sunday - Saturday: 6:00am - 7:00/8:00pm" in season, and on the home page, "Winter Hours Are Inconsistent, Leave a Voicemail." No second line and no answering service anywhere on the site. Brian is listed as "Owner & Captain of The Irish," which is why we guess he isn't at a desk much in August, though you'd know better.

We checked the others on the bay before calling any of this a gap, and mostly it isn't. You run FareHarbor with a Book Now button, and so do Emerald Air Service, Smokey Bay Air, Destination Alaska Charters and Bear Viewing Alaska, while Central Charters takes instant online booking through a listing site. On booking you are level with all of them, and anyone telling you to fix that system is wrong on the facts. Not one operator we checked in Homer or Kachemak Bay runs AI guest response of any kind. The one real difference is reach. Emerald Air Service publishes text and WhatsApp alongside phone and email. You publish a phone line, an inbox and voicemail.

One phone line, one shared inbox, no answering service

What your contact page publishes, and what it doesn't. We can only see what you publish, so treat this as the outside view of your setup rather than a description of it.

02

You already promised to resell cancelled seats, and the clock decides whether you can

Your own policy sets the stakes better than we could.

Your policy prices the run-up to a departure, and it does something more unusual too. You wrote that if you resell a cancelled space you don't keep the deposit.

That is a promise with a clock on it. Reselling is what lets you hand the deposit back, and reselling gets harder the later you learn the seat is free. Your own paragraph says you do everything you can to resell. The build is not about trying harder, it is about the message arriving sooner.

Thirty dollars a head

What you charge if you leave the dock and can't fish, from your FAQ. Weather before departure is a full refund. You carry both, so timing is your cost, not the guest's.

03

The change is that a message stops waiting for someone to open the inbox

One cancellation, walked step by step. We're guessing at your side of it, so correct us.

Nothing about how guests reach you changes. No new number, no new inbox, nothing to sign into.

A guest emails at 9pm to cancel Saturday

TODAY	AFTER
1 A guest emails at 9pm to cancel Saturday.	A guest emails at 9pm to cancel Saturday. Same channel, same message.
2 It lands in the shared inbox with everything else that evening.	The guest gets a written answer inside a minute, in your FAQ's own wording, asking them to call and confirm.
3 Nobody knows it is there until someone next opens the inbox.	One text reaches whoever is covering the inbox, up to six a day, after which it waits for the evening summary.
4 Whoever opens it first works out how close Saturday is.	A person opens FareHarbor and checks the date. Ten seconds.
5 Someone calls the guest back to confirm.	Someone calls the guest back to confirm.
6 However long that took is however long the seat sat unsellable.	The seat is known to be free while there is still time to act on it.

The Today column is our reading of how this probably goes, not something we watched, so strike anything we got wrong. Look at the fifth row. The call to the guest is identical on both sides, deliberately. That conversation is the part of your business worth protecting, so we're not putting a machine near it.

What the columns can't show is the discipline underneath. Nothing marks a message read, files it or hides it, and the evening summary lists everything that arrived. A system that pings you about the urgent ones can quietly train you to stop scanning the inbox, and then one miss is worse than never having had it. Read-only access makes that impossible rather than promised.

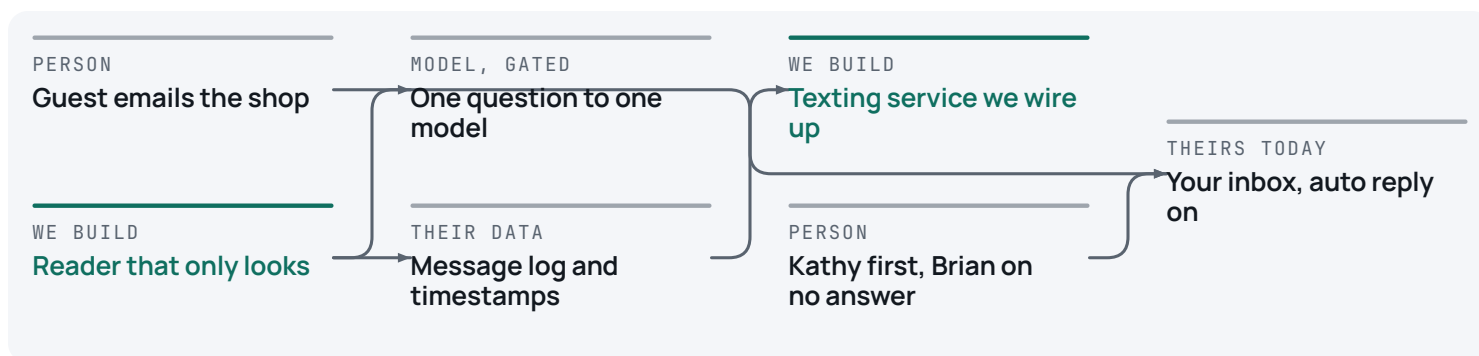
What we would build is small, and most of it isn't AI

One model call, one text message, and a count nobody has taken.

The moving pieces, in plain nouns. An automatic reply, a setting inside your own Gmail. A reader that watches new mail with a key that can only look, never change. One short question put to a language model, is this about cancelling. A log with a row per message. A text to whichever phone you name.

The automatic reply goes out in week one, before any model exists. Every message gets the same three lines. The middle one is your FAQ sentence character for character. The other two, a line saying it is automatic and nobody has read it yet plus your phone number, are ours and you sign them off. It says nothing about any booking, so it can't be wrong about one.

The reader copies each new message, blanks the sender's name, address and phone number by matching the header rather than by a second model, and asks one question of one model. Any day word the guest typed is lifted in by a fixed word list, with no interpretation and no check against your calendar. On a yes, one text goes to whoever you nominate, and if the guest still has no reply an hour later it goes to a second person. Six a day maximum, past which nothing more is sent and every message still appears on the evening summary.



One model call in the whole picture. Everything else is plumbing, arithmetic, or a person.

We looked hard at whether this needs AI at all, and mostly it doesn't. The date arithmetic belongs in code where it can be tested, not in a model where it can be invented. The guest reply is one canned sentence. Which refund rule applies on a

weather day is a two-way branch out of your own policy. Being reachable when nobody is at the desk is a phone company purchase.

One job resists a rule. A keyword filter catches "I need to cancel" and misses "family emergency, we won't make Saturday." People write it a hundred ways and the point is catching all of them. That is the only place we put a model.

Where we would not use it: anywhere a guest gets bad news. A controlled study of 467 people found that when a booking got worse, the case tested being a price increase, willingness to book dropped further if a chatbot delivered it than if a person did. Weather cancellations look like the same moment, which is our reading rather than something that study tested.

Almost all of it is bought. The reply is a Gmail setting, the mailbox access is Google's own, the model is off the shelf. Two things are ours, the small program in the middle and the labelled set of your real messages that shows whether the model is any good. That second one you keep.

One thing we won't build. FareHarbor, which you already pay for, has announced FareHarbor Agent, described on their own page as answering guest questions across phone, email, chat and social, rolling out to a select group of operators. Ask them for it. If it covers this, we stop billing.

05

The fortnight pays for itself only if you'd otherwise have bought the build blind

The honest return is a number you don't have, plus a decision you might not make.

There is no dollar return here for the bigger build. What it gets instead is a break-even you can test against your own season.

Your policy makes the arithmetic work. If a cancelled seat goes unsold you keep the fifty percent deposit. If you resell it you hand that deposit back and collect the full rate from the new guest, so reselling is worth about half a seat, not a whole one. Using the whole seat would have halved the threshold in our favour and it would have been wrong.

At 5,500 over three years the Clock costs about 10,900 with contingency, at 8,000 about 16,800. Half a seat at 250, 350 or 450 dollars, all three our guesses because you publish no rate, allowing a slow first year, gives between 19 and 54 recovered seats a season to break even. Cheapest build and dearest seat is 19, dearest build and cheapest seat is 54. Ask whether that many late cancellations even happen to you.

The measurement fortnight, plus the one assumption gate two will test later

	Conservative	Most likely	Aggressive
ⓘ Emails a year we assume reach you, roughly two to seven a day	750	1,400	2,500
ⓘ Cancellation emails we assume one model call catches	90 in 100	92 in 100	95 in 100
ⓘ Chance the fortnight tells you to build nothing	40%	55%	70%
☑ Your deposit at 14 days or less, from your FAQ	50%	50%	50%
☑ Total cost, our fee plus 8, 6 or 4 hours of your time	\$3,960	\$3,747	\$3,507
☑ Spend avoided plus time saved, credited in year one	\$4,110	\$7,720	\$13,718
☑ Share of cost recovered, counting the whole decision	104%	206%	391%
☑ Share recovered counting only year one of that build	83%	166%	323%

We valued an hour of your time at 50 dollars throughout. Our figure, not your payroll. Cost, conservative column: our fee 2,800 fixed, plus eight hours of your time at setup, plus two hours a year afterwards watching the automatic reply, times twenty percent contingency, gives 3,960. The middle column is six hours and eighteen percent, the right is four and fifteen. The fee never moves. Benefit, all credited in year one because that is when you decide. Building blind commits you to 9,500, being 5,500 to build plus 400 of your time plus 1,200 a year of running for three years. Times the 40 percent chance we tell you not to build, 3,800. Add 180, being a 1,200 duplicate-vendor exposure at a 15 percent chance. Add 480, being 800 of labelling at 60 percent. and 375 of first-touch time the reply saves that year, which is capacity rather than cash and only becomes money if that time goes to something that earns. Total 4,835, cut fifteen percent as a risk adjustment because not every line lands in full, gives 4,110. The other columns move the same drivers. That benefit is mostly money NOT SPENT, avoided once when you decide. Credit only the first year of the build and the conservative case recovers 83 percent, which doesn't clear, and it only counts at all if you'd otherwise have bought blind. The 40 percent is the first number to argue with, because at 20 the conservative case fails. Run the reply alone over three years with no decision value and it recovers about a quarter of its own cost. It is not what pays for this.

✓ verified against a source ☑ modelled from stated drivers

ⓘ assumed, and named as such

104%

of the fortnight's cost comes back in the conservative case, counting the whole three-year commitment you would avoid. Count only year one of that build and it is 83 percent, which doesn't clear. Both are in the table, and both depend on you having been willing to buy blind.

The outside view first, because it applies to us too. We could not find a single independently measured AI deployment at family-run charter scale anywhere, and the small-operator case studies we did find were published by the vendors selling the product. The closest sector the national survey tracks to yours sits around 8 percent adoption, near the bottom. The reason this one is built to survive that record is that the only thing funded today is a count, and the count is allowed to end it.

Two assumptions above become actuals from your own records at the end of the fortnight, emails a day and how many are cancellations. The 40 percent never becomes

an actual, it just resolves. We'd propose Kathy owns the number, since she is co-owner and office manager by your own page.

06

Start with two weeks of counting, and keep the exit real

Sequenced so the cheapest, most reversible thing happens first.

Your season winds down through September, we assume, so a fortnight starting late August would measure real summer traffic. If that is wrong, say so and we'll move it. The Clock and everything depending on it run next season, because gate three needs a full in-season month and there won't be one left. The off-season desk is the one Next item that could run this winter.

Phase	What runs	How we would know it worked
Now	Ask FareHarbor for the Agent pilot. Free, and we recommend it whether or not you hire us.	A written yes or no.
	Switch on the automatic reply, your FAQ's own wording, no AI in it.	Share of inbound email acknowledged, against whatever it is now.
	Sign the data agreement and grant read-only mailbox access.	All three dated before the first logged message.
	Two weeks of logging and labelling every inbound email. The whole 2,800 dollar ask.	Emails a day, how many are about cancelling, and median minutes to a human seeing them.
	One call to your carrier to find what that line offers.	A written answer on transcription.
	The written verdict, before your season winds down.	A yes or no against a threshold both sides signed on day one.
Next	The Clock itself, email only. Indicative 5,500 to 8,000 dollars, month to month, released only if the count clears your threshold.	Median minutes from a cancellation email to a human confirming by phone.
	A read-only FareHarbor lookup so an alert can carry days to departure.	Share of alerts carrying a correct date.

Four gates, and the money moves one at a time. The count that decides gate one is signed before anything starts, so it can't be rationalised backwards once the data lands. Gate one is the real exit. If the fortnight comes in under that bar, our recommendation says the build is not worth making, you keep the automatic reply and the numbers, and you buy nothing further. On today's evidence that outcome is genuinely live. Gate two, the measured catch rate has to reach the ninety in a hundred assumed above. Gate three, the Clock has to beat the baseline over a full in-season month. Gate four, nothing later is funded until the number it depends on exists.

Phase	What runs	How we would know it worked
	Re-measure in the May to June shoulder. The Clock is released on an August-only count, which we think fair for a build this cancellable, and nothing bigger should rest on it.	How far the shoulder baseline moves.
	The off-season desk, an autoresponder and an interest form. No chatbot, no model.	Off-season inquiries captured.
	Weather-day notifications, templates with a human on the send button.	Minutes to every party told, zero rule errors.
Later	Filling a freed seat from a standby list. A database query and a text, no AI in it.	Consented names, then seats resold before departure.
	Measuring FareHarbor Agent against the same baseline if the pilot lands.	The same numbers before and after.
	Public review replies, a permanent do-not-automate zone. An investigation found AI-written review summaries softened serious complaints, and a name on the Spit since 1979 is the asset that damages.	AI-written review replies, zero, permanently.

Four gates, and the money moves one at a time. The count that decides gate one is signed before anything starts, so it can't be rationalised backwards once the data lands. Gate one is the real exit. If the fortnight comes in under that bar, our recommendation says the build is not worth making, you keep the automatic reply and the numbers, and you buy nothing further. On today's evidence that outcome is genuinely live. Gate two, the measured catch rate has to reach the ninety in a hundred assumed above. Gate three, the Clock has to beat the baseline over a full in-season month. Gate four, nothing later is funded until the number it depends on exists.

A signed data agreement and consent from both of you, before any access. Read-only mailbox access through Google's own consent screen, revocable in one click, no password ever shared. A forwarding rule works if you'd rather grant nothing, and it costs you two things, it shows what arrives but never when someone answered, so there is no median-minutes number and no second alert, and gate three then has no baseline.

About eight hours of your time across the fortnight, priced into the table rather than called included. The scheduled part is small, fifteen minutes each on day one, signing off the reply wording, one call to your carrier and ten minutes a week checking labels. The rest is padding for interruptions, and we'd rather over-price your time than under-price it.

07

Four things that could make this the wrong call, and what we'd do about each

Every one of these is a real reason to say no.

Your existing vendor may do this for free. FareHarbor Agent already covers email, so ask them for the pilot before you spend anything with us, and if it covers this we stop billing.

The ceiling may be small. If two weeks turn up a handful of cancellation emails the right answer is that you build nothing, which is why the threshold is agreed before any number exists.

Late August only tells you about late August, so the finding gets labelled August-only rather than folded into something bigger.

We may be reading one sentence too hard. That FAQ line could be old policy text that no longer describes today, and two weeks of real messages tests it. "The channel is fine" is a legitimate result and we would write it.

On the word biggest, which appears twice. FareHarbor is the biggest risk to whether you hire us at all. The 40 percent is the biggest driver of whether the arithmetic works.

The next step is a reply, not a call

If this was useful, tell us. If it was wrong somewhere, tell us that instead, and there are places above where we guessed at your side of things. Nothing to book and nothing to sit through. If you want the fortnight, the sooner it starts the more of your season it sees.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1 Your hours and the winter voicemail line
<https://www.northcountrycharters.com/>
- 2 The call-and-confirm request and your cancellation and weather rules
<https://www.northcountrycharters.com/faq>
- 3 Brian Ritchie and Kathy Brennan, and their published titles
<https://www.northcountrycharters.com/names-you-might-know>
- 4 One phone number and one shared inbox
<https://www.northcountrycharters.com/contact-1>
- 5 FareHarbor Agent and its select rollout
<https://fareharbor.com/blog/meet-fareharbor-agent-ai-powered-guest-communication-for-tour-operators/>
- 6 Emerald Air Service on text, WhatsApp and a 2027 portal
<https://www.emeraldairservice.com/>
- 7 Central Charters taking instant online booking
<https://fishingbooker.com/charters/view/41714>
- 8 Smokey Bay Air on FareHarbor
<https://smokeybayair.com/bear-viewing.html>

- 9 Destination Alaska Charters on FareHarbor
<https://destinationalaskacharters.com/>
- 10 Bear Viewing Alaska on FareHarbor
<https://www.bearviewingalaska.com/>
- 11 The 467-participant study on a chatbot versus a person delivering a price increase
<https://www.mdpi.com/2673-5768/6/1/36>
- 12 Your industry near the bottom of AI adoption
<https://www.fool.com/research/ai-adoption-statistics/>
- 13 AI review summaries downplaying complaints
<https://travel.yahoo.com/advice/safety/articles/booked-disaster-tripadvisor-ai-accused-223127463.html>