

FIELD STUDY PREPARED FOR NORTH COUNTRY CHARTERS

Your FAQ asks guests to phone because voicemails and emails aren't always seen, and real money rides on that.

HOMER, ALASKA 2 AUGUST 2026

IN SHORT

We spent this week studying North Country Charters. Nobody asked us to. We do this for one Alaska business at a time, we give it away, and this is yours.

What we looked at: everything you publish. Your home page, your FAQ, your captains, your contact details. Then your neighbours on the Spit, Central Charters, Smokey Bay Air, Emerald Air Service, Destination Alaska Charters, Bear Viewing Alaska, and what the booking platforms behind them actually do. Then the honest research on what AI is doing for operators your size, which is less than the people selling it suggest.

What we found sits in one sentence you wrote yourselves. Your FAQ says: "IF YOU NEED TO CANCEL please call us and confirm over the phone with us that you are cancelling. We don't always see voicemails or emails right away and every minute counts." Directly beneath it, your own policy attaches money to timing. Fifty percent at fourteen days or less. One hundred percent at twenty four hours or a no-show. You have written down that a channel is unreliable, and you have written down what the clock is worth.

Be clear about what that is and is not. It is a fact about a message channel. It is not us telling you that you're overwhelmed, and we looked for that. Nothing you publish says the phones ring off the hook, so we won't imply it.

What we'd build first has no AI in it at all. A plain automatic reply on your shared inbox, in the wording already on your own FAQ, so every guest who writes in gets your call-and-confirm request inside a minute instead of only the ones who read the FAQ page. That is a Gmail setting and about ten minutes of work.

What we'd ask you to buy is two weeks of counting. Two thousand eight hundred dollars, fixed. Nothing logs how many messages reach you or how many are about cancelling, so no honest price exists for the real build. The fortnight produces that number, and it ends with a written verdict that includes "build nothing further."

What we couldn't verify: your charter rates, your message volume, and whether that FAQ line still describes today. So there is no dollar return in here for the bigger build, only a break-even you can test yourself.

The ask is a reply. Not a call.

01

You already wrote down where this hurts, on your own FAQ

The strongest thing we found about your business, you published yourselves.

Most studies like this open with a problem the writer invented. Yours was already on the page. Your FAQ tells guests, in capitals, to phone and confirm a cancellation rather than trust the message they just sent, and it says why in the same breath.

Read that against the rest of what you publish. One number. One shared inbox. Hours of "Sunday - Saturday: 6:00am - 7:00/8:00pm" in season, and on the home page,

"Winter Hours Are Inconsistent, Leave a Voicemail." No second line and no answering service anywhere on the site. Brian is listed as "Owner & Captain of The Irish," so for much of a summer day he is on the water.

One phone line, one shared inbox, no answering service

What your contact page publishes, and what it doesn't. Outside your posted hours the only path in is voicemail.

One thing we want to be careful about. We are not saying you're drowning, and we went looking. No page of yours says the phones ring off the hook or that calls get missed. What is there is a written acknowledgement that one channel is slow, sitting above the paragraphs where the money changes. That is a smaller claim than a salesperson would make, and it is the one the evidence supports.

02

The clock costs you in two directions, and your own policy sets both

Your cancellation rules are the clearest statement of what timing is worth here.

A charter seat is perishable in a way a hotel room barely is. Once the boat leaves, that seat earned nothing and can never be sold again. Your policy already prices the run-up to that moment. A cancellation that reaches you late doesn't change what the guest owes, since they gave notice when they sent it. What it changes is how long that seat had to find someone else.

Two weeks, then one day

The two lines in your published policy where the money changes. Both count from the guest's notice, not from when it was read.

There is a second direction, and it ends up in reviews. A guest who emailed four days out, then got charged as though they hadn't, is holding written proof they told you. Nobody wins that conversation.

We can't tell you what either costs you a season. That number doesn't exist, and inventing one would be the fastest way to lose your trust.

03

The change is that a message stops waiting for someone to open the inbox

One real cancellation, walked step by step, before and after.

Nothing about how guests reach you changes. No new number, no new inbox, no app, nothing to sign into. The work happens around the channel you already have.

A guest emails at 9pm to cancel Saturday

TODAY	AFTER
1 A guest emails at 9pm to cancel Saturday.	A guest emails at 9pm to cancel Saturday. Same channel, same message.
2 It lands in the shared inbox with everything else that evening.	The guest gets a written answer inside a minute, in your FAQ's own wording, asking them to call and confirm.
3 Nobody knows it is there until someone next opens the inbox.	One text reaches Kathy's phone saying a message reads like a cancellation.
4 Whoever opens it first works out how close Saturday is.	Kathy opens FareHarbor and checks the date herself. Ten seconds.
5 Kathy calls the guest back to confirm.	Kathy calls the guest back to confirm.
6 However long that took is however long the seat sat unsellable.	The seat is known to be free while there is still time to act on it.

Look at the fifth row. The call to the guest is identical on both sides, deliberately. That conversation is the part of your business worth protecting, so we're not putting a machine near it. The fourth row barely moves either, because checking how close Saturday is means opening FareHarbor, and we'd rather a person did that than have software guess at a date that governs real money.

What the columns can't show is the discipline underneath. Nothing marks a message read, files it, or hides it, and the evening summary lists everything that arrived. A system that pings you about the urgent ones can quietly train you to stop scanning the inbox, and then one miss is worse than never having had it. Read-only access makes that impossible rather than promised.


What we would build is small, and most of it isn't AI

One model call, one text message, and a count nobody has taken.

The moving pieces, in plain nouns. An automatic reply, which is a setting inside your own Gmail. A reader that watches new mail with a key that can only look, never change. One short question put to a language model, is this message about cancelling. A log with a row per message. A text to a phone Kathy already carries.

The automatic reply goes out first, in week one, before any model exists. Every message gets the same fixed sentence from your FAQ, so it can't say anything wrong about a booking, because it says nothing about any booking. You approve the wording first.

The reader copies each new message, blanks out the name and phone number, and asks one question of one model. When the answer is yes, one text goes to Kathy. If the guest still has no reply an hour later, one goes to Brian. Six a day maximum, and a word you can text back that silences it.



One model call in the whole picture. Everything else is plumbing, arithmetic, or a person.

We looked hard at whether this needs AI at all, and mostly it doesn't. The date arithmetic belongs in code where it can be tested, not in a model where it can be invented. The guest reply is one canned sentence. Which refund rule applies on a weather day is a two-way branch out of your own policy. Being reachable while a captain is on the water is a phone company purchase, not software.

One job resists a rule. A keyword filter catches "I need to cancel" and misses "family emergency, we won't make Saturday." People write it a hundred ways and the point is catching all of them. That is the only place we put a model.

Where we would not use it: anywhere a guest gets bad news. A study of 467 people found willingness to book fell further with an AI agent than with a human when the news was bad. Weather cancellations are that moment for you.

Almost all of it is bought. The automatic reply is a Gmail setting, the mailbox access is Google's own, the model is off the shelf. Two things are ours to build, the small program in the middle, and the labelled set of your real messages that shows whether the model is any good. That second one you keep.

One thing we won't build. FareHarbor, which you already pay for, has announced FareHarbor Agent, described on their own page as answering guest questions across phone, email, chat and social, and rolling out to a select group of operators. Ask them for it. If it covers this, we stop billing.

05

We're asking for 2,800 dollars, and the conservative case clears it by four points

The honest return of the fortnight is a number you don't have, not a saving we projected.

There is no dollar return here for the bigger build, because producing one would mean inventing your charter rate and your message volume. Neither is published anywhere and we won't make them up.

What it gets instead is a break-even you can test against your own season. At our indicative 5,500 to 8,000 dollars plus running costs over three years, the Clock pays for itself somewhere between 19 and 54 recovered seats a season. At 350 dollars a seat, our assumption and not your number, that is 24 seats at the low end of our price and 38

at the high end. We count only half a seat's value, because your own deposit already protects the other half.

Ask whether that many late cancellations even happen to you. If the answer is obviously not, you have your verdict, and it cost you nothing.

The measurement fortnight, the only thing being asked for today

	Conservative	Most likely	Aggressive
Emails a year we assume reach your inbox	750	1,400	2,500
Chance the fortnight says build nothing	40%	55%	70%
Your deposit at 14 days or less, from your FAQ	50%	50%	50%
Total cost, our fee plus your time	\$3,960	\$3,747	\$3,507
Value back in year one	\$4,110	\$7,720	\$13,718
Share of cost recovered	104%	206%	391%
Payback	12 months	6 months	4 months

Assumptions, all ours and all arguable. Our fee is fixed and can't overrun, and we still added contingency on top. Your time is priced at an assumed 50 dollars an hour, our figure, not your payroll. The largest driver is the chance the count says build nothing, set at 40 percent in the conservative case. That is the first number you should argue with, because if you think it should be 20 percent the conservative case fails and you shouldn't buy the fortnight. One result we'd rather show than bury: the autoresponder alone recovers 25 percent of the fee over three years. It is not what pays for this. The number is.

verified against a source

modelled from stated drivers

assumed, and named as such

104%

of the fortnight's cost comes back in the conservative case. It clears by four points, not by a landslide, and it clears inside one year rather than the five most cases are stretched over.

The outside view first, because it applies to us too. Roughly 95 percent of enterprise AI pilots produce no measurable impact on the books. Your industry sits near the bottom of AI adoption, around 8 percent, and we found no independently measured deployment at family-run charter scale anywhere. That absence is not a reason to be first. It is why we're selling two weeks of counting instead of a build.

Kathy Brennan owns the number, because it lands in the inbox she runs and every alert routes to her first. At the end of the fortnight three assumptions above get replaced with actuals from your own records.

06

Start with two weeks of counting, and keep the exit real

Sequenced so the cheapest, most reversible thing happens first.

One date matters most. Your season closes mid-September, so a fortnight started by 24 August measures real summer traffic. Miss it and the next honest window is May 2027. The off-season work runs the other way and gets better as the season closes.

Phase	What runs	How we would know it worked
Now	Ask FareHarbor for the Agent pilot. Costs nothing, and we recommend it whether or not you hire us.	A written yes or no, plus a start date.
	Switch on the automatic reply, in your FAQ's own wording, with no AI in it.	Every email answered inside a minute, from a baseline of none.
	Sign the data agreement and grant read-only mailbox access.	All three dated before the first logged message.
	Two weeks of logging and labelling every inbound email. This is the whole 2,800 dollar ask.	Emails a day, how many are about cancelling, and median minutes to a human seeing them.
	One call to your carrier to find what that line offers.	A written answer on whether transcription exists.
	The written verdict, delivered inside your season.	A yes or no against a threshold both sides signed before day one.
Next	The Clock itself, email only. Indicative 5,500 to 8,000 dollars, month to month, released only if the count clears your threshold.	Median minutes from a cancellation email to a human confirming by phone.
	A read-only FareHarbor lookup so an alert can carry days to departure.	Share of alerts carrying a correct date stamp.

Four gates, and the money moves one at a time. The count that decides gate one is written down and signed before anything starts, so it can't be rationalised backwards once the data lands. Gate one is the real exit. If the fortnight comes in under that bar, our recommendation says the build is not worth making, you keep the automatic reply and the numbers, and you buy nothing further from us. On today's evidence that outcome is genuinely live. Gate two, measured accuracy has to land inside the range we assumed. Gate three, the Clock has to beat the baseline over a full in-season month. Gate four, nothing later is funded until the number it depends on exists.

Phase	What runs	How we would know it worked
	Re-measure in the May to June shoulder before anything rests on August numbers.	How far the shoulder baseline moves.
	The off-season desk, an autoresponder and an interest form. No chatbot, no model.	Off-season inquiries captured, against a voicemail baseline of nothing.
	Weather-day notifications, templates with a human on the send button.	Minutes to every party told, and zero refund-rule errors.
Later	Filling a freed seat from a standby list. A database query and a text, with no AI in it.	Consented names on the list, zero today.
	A staffable text path. Emerald Air Service publishes text and WhatsApp. We are not the vendor for this one.	Share of after-hours inquiries answered before morning.
	The Guest Room attach question, which starts as a measurement and may never become a build.	Room nights attached per hundred bookings.
	Measuring FareHarbor Agent against the same baseline if the pilot lands.	The same numbers before and after it is switched on.
	Public review replies, a permanent do-not-automate zone.	AI-written review replies, zero, permanently.

Four gates, and the money moves one at a time. The count that decides gate one is written down and signed before anything starts, so it can't be rationalised backwards once the data lands. Gate one is the real exit. If the fortnight comes in under that bar, our recommendation says the build is not worth making, you keep the automatic reply and the numbers, and you buy nothing further from us. On today's evidence that outcome is genuinely live. Gate two, measured accuracy has to land inside the range we assumed. Gate three, the Clock has to beat the baseline over a full in-season month. Gate four, nothing later is funded until the number it depends on exists.

A signed data agreement and consent from both of you, before any access. Read-only mailbox access through Google's own consent screen, revocable in one click, and no password ever shared. A forwarding rule works instead if you'd rather. Fifteen minutes from both of you on day one to agree the threshold and what counts as a cancellation. Your sign-off on the reply wording. Ten minutes a week from Kathy checking our labels.

07

Four things that could make this the wrong call, and what we'd do about each

Every one of these is a real reason to say no.

Your existing vendor may do this for free. This is the biggest commercial risk here and we haven't retired it. FareHarbor Agent already covers email, and alerting an operator is an easier feature for them to add than the guest-facing one they've shipped. So we tell you to ask for the pilot now, before you spend anything with us.

The honest ceiling may be small. If two weeks turn up a handful of cancellation emails, the right answer is that you should build nothing. That exit is written in with a threshold agreed before any number exists, because a gate nobody would honour is worse than no gate. You keep the reply, you keep the count, and the money bought the number that stopped a larger mistake.

Late August only tells you about late August, so the finding gets labelled August-only rather than folded into a build to make the annual case look bigger.

We may be reading one sentence too hard. That FAQ line could be old policy text that no longer describes how anything works today. Two weeks of real messages tests it, and "the channel is fine" is a legitimate result that ends with us recommending you build nothing.

The next step is a reply, not a call

If this was useful, tell us. If it was wrong somewhere, tell us that instead, we'd rather know. There is nothing to book and nothing to sit through. If you want the fortnight, we'd need access by 17 August to start on the 24th.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1 Your hours, the winter voicemail line, and since 1979
<https://www.northcountrycharters.com/>
- 2 The call-and-confirm request and your cancellation and weather rules
<https://www.northcountrycharters.com/faq>
- 3 Brian Ritchie and Kathy Brennan, and their published titles
<https://www.northcountrycharters.com/names-you-might-know>
- 4 One phone number and one shared inbox
<https://www.northcountrycharters.com/contact-1>
- 5 Local, Woman Co-Owned Fishing in Homer AK
<https://www.northcountrycharters.com/about>
- 6 FareHarbor Agent and its select rollout
<https://fareharbor.com/blog/meet-fareharbor-agent-ai-powered-guest-communication-for-tour-operators/>
- 7 Emerald Air Service on text, WhatsApp and a 2027 portal
<https://www.emeraldairservice.com/>
- 8 Central Charters taking instant online booking
<https://fishingbooker.com/charters/view/41714>

- 9 Smokey Bay Air on FareHarbor
<https://smokeybayair.com/bear-viewing.html>
- 10 Destination Alaska Charters on FareHarbor
<https://destinationalaskacharters.com/>
- 11 Bear Viewing Alaska on FareHarbor
<https://www.bearviewingalaska.com/>
- 12 The 467-participant study on AI versus a human on bad news
<https://www.mdpi.com/2673-5768/6/1/36>
- 13 Your industry near the bottom of AI adoption
<https://www.fool.com/research/ai-adoption-statistics/>
- 14 AI review summaries downplaying complaints
<https://travel.yahoo.com/advice/safety/articles/booked-disaster-tripadvisor-ai-accused-223127463.html>
- 15 The one small-operator case study, published by its vendor
<https://www.vendasta.com/blog/call-answering-service-for-small-business/>
- 16 North Country listed independently of your own site
<https://www.fishalaskamagazine.com/outfitters/north-country-charters/>