

FIELD STUDY PREPARED FOR NORTH COUNTRY CHARTERS

Your FAQ asks guests to phone because voicemails and emails aren't always seen, and real money rides on that.

HOMER, ALASKA 2 AUGUST 2026

IN SHORT

Your FAQ carries one paragraph that does two things at once. It opens with "We have phenomenal office staff and we do everything we can to resell cancelled seats," and it ends with "IF YOU NEED TO CANCEL please call us and confirm over the phone with us that you are cancelling. We don't always see voicemails or emails right away and every minute counts."

Both halves are true together, and the second is not a complaint about the first. Your office can be as good as you say and a message can still sit unseen, because those are different problems. One is about people, the other is about a channel, and everything below is about the channel. We are not telling you that you're overwhelmed. Nothing you publish says the phones ring off the hook, so we won't imply it.

We spent this week on your business, unasked and at no charge. What we read: everything you publish, then other Homer operators, Central Charters, Smokey Bay Air, Emerald Air Service, Destination Alaska Charters, Bear Viewing Alaska, then the

honest research on what AI does for operators your size, which is less than the people selling it suggest.

What we'd build first has no AI in it. A plain automatic reply on your shared inbox carrying your own FAQ sentence, so every guest who writes in gets your call-and-confirm request instead of only the ones who read the FAQ page. That is a Gmail setting and ten minutes of work.

What we'd ask you to buy is two weeks of counting. Two thousand eight hundred dollars fixed, plus roughly eight hours of your own time. Nothing logs how many messages reach you or how many are about cancelling, so no honest price exists for the real build. The fortnight produces that number, and it ends with a verdict that includes "build nothing further."

On the money, the conservative case gets about all of that back inside a year. Read why before you believe it. The largest reason is our own estimate that there is a four in ten chance we tell you not to build, and that is a number you should argue with.

What we couldn't verify: your charter rates, your message volume, and whether that FAQ line still describes today. So there is no dollar return here for the bigger build, only a break-even you can test yourself.

The ask is a reply. Not a call.

01

You already wrote down where this hurts, and you wrote the answer to it too

The strongest thing we found about your business, you published yourselves.

Your FAQ sets out your money lines first, fifty percent at fourteen days or less, the full rate at twenty four hours or a no-show. It then says something most operators would

not put in writing: "If we can resell cancelled spaces, then we do not retain the deposit." It closes by telling guests, in capitals, to phone and confirm a cancellation rather than trust the message they just sent.

Read that against the rest of what you publish. One number. One shared inbox. Hours of "Sunday - Saturday: 6:00am - 7:00/8:00pm" in season, and on the home page, "Winter Hours Are Inconsistent, Leave a Voicemail." Your site publishes no second line and no answering service. Brian is listed as "Owner & Captain of The Irish," which is our reason for guessing he isn't at a desk much in August, though you'd know better.

One phone line, one shared inbox, no answering service

What your contact page publishes, and what it doesn't. We can only see what you publish, so treat this as the outside view of your setup rather than a description of it.

One thing we want to be careful about. We are not saying you're drowning, and we went looking. No page of yours says the phones ring off the hook, and the same paragraph that admits the channel is slow also says your office staff are phenomenal. We take both at face value. A good office and a slow channel are different things, and only one of them is a software problem.

02

You already promised to resell cancelled seats, and the clock decides whether you can

Your own policy sets the stakes better than we could.

A charter seat is perishable in a way a hotel room barely is. Once the boat leaves, that seat earned nothing and can never be sold again. Your policy prices the run-up to that moment, and it does something more unusual too. You wrote that if you resell a cancelled space you don't keep the deposit.

That is a promise with a clock on it. Reselling is what lets you hand the deposit back, and reselling gets harder the later you learn the seat is free.

Thirty dollars a head

What you charge if you leave the dock and can't fish, from your FAQ. Weather before departure is a full refund. You carry both, so timing is your cost, not the guest's.

Your own paragraph says you do everything you can to resell. The build is not about trying harder, it is about the message arriving sooner.

We can't tell you what any of this costs you a season. That number doesn't exist, and inventing one would be the fastest way to lose your trust.

03

The change is that a message stops waiting for someone to open the inbox

One cancellation, walked step by step. We're guessing at your side of it, so correct us.

Nothing about how guests reach you changes. No new number, no new inbox, no app, nothing to sign into. The work happens around the channel you already have.

A guest emails at 9pm to cancel Saturday

TODAY	AFTER
1 A guest emails at 9pm to cancel Saturday.	A guest emails at 9pm to cancel Saturday. Same channel, same message.
2 It lands in the shared inbox with everything else that evening.	The guest gets a written answer inside a minute, in your FAQ's own wording, asking them to call and confirm.
3 Nobody knows it is there until someone next opens the inbox.	One text reaches whoever is covering the inbox, saying a message reads like a cancellation.
4 Whoever opens it first works out how close Saturday is.	A person opens FareHarbor and checks the date. Ten seconds.
5 Someone calls the guest back to confirm.	Someone calls the guest back to confirm.
6 However long that took is however long the seat sat unsellable.	The seat is known to be free while there is still time to act on it.

The Today column is our reading of how this probably goes, not something we watched, so strike anything we got wrong. Look at the fifth row. The call to the guest is identical on both sides, deliberately. That conversation is the part of your business worth protecting, so we're not putting a machine near it. The fourth row barely moves either, because checking how close Saturday is means opening FareHarbor, and we'd rather a person did that than have software guess at a date that governs real money.

What the columns can't show is the discipline underneath. Nothing marks a message read, files it, or hides it, and the evening summary lists everything that arrived. A system that pings you about the urgent ones can quietly train you to stop scanning the inbox, and then one miss is worse than never having had it. Read-only access makes that impossible rather than promised.


What we would build is small, and most of it isn't AI

One model call, one text message, and a count nobody has taken.

The moving pieces, in plain nouns. An automatic reply, which is a setting inside your own Gmail. A reader that watches new mail with a key that can only look, never change. One short question put to a language model, is this message about cancelling. A log with a row per message. A text to whichever phone you tell us to send it to.

The automatic reply goes out first, in week one, before any model exists. Every message gets the same three lines. The middle one is your own FAQ sentence character for character. The other two, a line saying it is automatic and nobody has read the message yet plus your phone number, are ours and you sign them off. It says nothing about any particular booking, so it can't be wrong about one.

The reader copies each new message, blanks the sender's name, address and phone number by matching the header rather than by a second model, and asks one question of one model. Any day or month word the guest typed is lifted into the alert by a fixed word list, with no interpretation and no check against your calendar. When the answer is yes, one text goes to whoever you nominate, and if the guest still has no reply an hour later it goes to a second person. Six texts a day maximum, past which nothing more is sent and every message still appears on the evening summary. That is why the summary lists everything rather than the leftovers.



One model call in the whole picture. Everything else is plumbing, arithmetic, or a person.

We looked hard at whether this needs AI at all, and mostly it doesn't. The date arithmetic belongs in code where it can be tested, not in a model where it can be

invented. The guest reply is one canned sentence. Which refund rule applies on a weather day is a two-way branch out of your own policy. Being reachable when nobody is at the desk is a phone company purchase, not software.

One job resists a rule. A keyword filter catches "I need to cancel" and misses "family emergency, we won't make Saturday." People write it a hundred ways and the point is catching all of them. That is the only place we put a model.

Where we would not use it: anywhere a guest gets bad news. A controlled study of 467 people found that when a booking got worse, and the case tested was a price increase, willingness to book dropped further if a chatbot delivered the news than if a person did. Weather cancellations look like the same shape of moment, which is our reading rather than something that study tested.

Almost all of it is bought. The automatic reply is a Gmail setting, the mailbox access is Google's own, the model is off the shelf. Two things are ours to build, the small program in the middle, and the labelled set of your real messages that shows whether the model is any good. That second one you keep.

One thing we won't build. FareHarbor, which you already pay for, has announced FareHarbor Agent, described on their own page as answering guest questions across phone, email, chat and social, and rolling out to a select group of operators. Ask them for it. If it covers this, we stop billing.

05

We're asking for 2,800 dollars, and the conservative case clears it by four points

The honest return of the fortnight is a number you don't have, not a saving we projected.

There is no dollar return here for the bigger build, because producing one would mean inventing your charter rate and your message volume. Neither is published anywhere and we won't make them up.

What it gets instead is a break-even you can test against your own season. Your policy is what makes the arithmetic work. If a cancelled seat goes unsold you keep the fifty percent deposit. If you resell it you hand that deposit back and collect the full rate from the new guest, so reselling is worth about half a seat to you, not a whole one. We used the half. Using the whole seat would have halved the threshold in our favour and it would have been wrong.

At our indicative 5,500 to 8,000 dollars plus running costs over three years, the Clock needs between 19 and 54 recovered seats a season to pay for itself. That spread covers both ends of our price and three different guesses at what a seat is worth, because we don't know what you charge.

Ask whether that many late cancellations even happen to you. If the answer is obviously not, you have your verdict, and it cost you nothing.

The measurement fortnight, the only thing being asked for today

	Conservative	Most likely	Aggressive
☐ Emails a year we assume reach your inbox	750	1,400	2,500
☐ Chance the fortnight says build nothing	40%	55%	70%
☐ Cancellation emails we assume one model call catches	90 in 100	92 in 100	95 in 100
☒ Your deposit at 14 days or less, from your FAQ	50%	50%	50%
☒ Total cost, our fee plus your time	\$3,960	\$3,747	\$3,507
☒ Spend avoided plus time saved, over the whole decision	\$4,110	\$7,720	\$13,718
☒ Share of cost recovered	104%	206%	391%
☒ Payback	12 months	6 months	4 months

How the benefit row is built, so you can rebuild it. Conservative case: 9,500 dollars is what building blind commits you to in total, 5,500 to build plus 400 of your time plus 1,200 a year of running for three years. Multiply by the 40 percent chance we tell you not to build, 3,800. Add 180 for the chance FareHarbor ships this and you avoid paying us twice, 480 for labelling you would otherwise buy inside the build, and 375 a year of first-touch time the autoresponder saves. That totals 4,835, and a partial first year brings it to 4,110 against 3,960 of cost. Read what that row is. It is mostly money NOT SPENT on a three-year commitment, avoided once when you decide, not cash arriving. It only counts if you would otherwise have bought the build without counting first. If you would never buy blind, the fortnight's return is the autoresponder and the number, nothing more. Assumptions, all ours and all arguable. The fee is fixed and can't overrun, and we still added contingency. Your time is assumed at eight hours priced at 50 dollars an hour, our figure, not your payroll. The biggest driver is that 40 percent, and it is the first number you should argue with, because at 20 the conservative case fails and you shouldn't buy the fortnight. One result we'd rather show than bury: the autoresponder alone recovers 25 percent of the fee over three years. It is not what pays for this. The number is.

✓ verified against a source ☑ modelled from stated drivers

ⓘ assumed, and named as such

104%

of the fortnight's cost comes back in the conservative case. It clears by four points, not by a landslide, and most of that is a three-year commitment avoided rather than cash arriving.

The outside view first, because it applies to us too. The closest sector the national survey tracks to yours, accommodation and food services, sits around 8 percent AI adoption, near the bottom, and we could not find a single independently measured AI deployment at family-run charter scale anywhere. The one small-operator case study we found was published by the vendor selling the product, in a different industry, which is not evidence and we won't dress it up as any. The reason this one is built to survive that record is that the only thing being funded today is a count, and the count is allowed to kill it.

We'd propose one named owner for the number rather than assume one. Kathy is the co-owner and office manager by your own page, so she is our suggestion, and you should override us if the inbox actually sits with someone else. At the end of the fortnight three assumptions above get replaced with actuals from your own records.

06

Start with two weeks of counting, and keep the exit real

Sequenced so the cheapest, most reversible thing happens first.

One question decides the timing and only you can answer it. We're assuming your season winds down through September, which means a fortnight starting in late August would measure real summer traffic. If that is wrong, say so and we'll move it.

Be clear about what fits. The fortnight and its verdict fit this season. Everything in the Next lane runs next season, because gate three needs a full in-season month and there won't be one left. The only decision actually on your table this year is whether to count.

Phase	What runs	How we would know it worked
Now	Ask FareHarbor for the Agent pilot. Costs nothing, and we recommend it whether or not you hire us.	A written yes or no, plus a start date.
	Switch on the automatic reply, in your FAQ's own wording, with no AI in it.	Share of inbound email getting a written acknowledgement, measured against whatever it is now, which we haven't seen.
	Sign the data agreement and grant read-only mailbox access.	All three dated before the first logged message.
	Two weeks of logging and labelling every inbound email. This is the whole 2,800 dollar ask.	Emails a day, how many are about cancelling, and median minutes to a human seeing them.
	One call to your carrier to find what that line offers.	A written answer on whether transcription exists.
	The written verdict, delivered before your season winds down.	A yes or no against a threshold both sides signed before day one.
Next	The Clock itself, email only. Indicative 5,500 to 8,000 dollars, month to month, released only if the count clears your threshold.	Median minutes from a cancellation email to a human confirming by phone.

Four gates, and the money moves one at a time. The count that decides gate one is written down and signed before anything starts, so it can't be rationalised backwards once the data lands. Gate one is the real exit. If the fortnight comes in under that bar, our recommendation says the build is not worth making, you keep the automatic reply and the numbers, and you buy nothing further from us. On today's evidence that outcome is genuinely live. Gate two, the measured catch rate has to reach the ninety in a hundred we assumed in the table above, tested against your own labelled messages. Gate three, the Clock has to beat the baseline over a full in-season month. Gate four, nothing later is funded until the number it depends on exists.

Phase	What runs	How we would know it worked
	A read-only FareHarbor lookup so an alert can carry days to departure.	Share of alerts carrying a correct date stamp.
	Re-measure in the May to June shoulder before anything rests on August numbers.	How far the shoulder baseline moves.
	The off-season desk, an autoresponder and an interest form. No chatbot, no model.	Off-season inquiries captured with a contact and a date, against a baseline the fortnight establishes.
	Weather-day notifications, templates with a human on the send button.	Minutes to every party told, and zero refund-rule errors.
Later	Filling a freed seat from a standby list. A database query and a text, with no AI in it.	Consented names on the list, then seats resold before departure.
	Measuring FareHarbor Agent against the same baseline if the pilot lands.	The same numbers before and after it is switched on.
	Public review replies, a permanent do-not-automate zone.	AI-written review replies, zero, permanently.

Four gates, and the money moves one at a time. The count that decides gate one is written down and signed before anything starts, so it can't be rationalised backwards once the data lands. Gate one is the real exit. If the fortnight comes in under that bar, our recommendation says the build is not worth making, you keep the automatic reply and the numbers, and you buy nothing further from us. On today's evidence that outcome is genuinely live. Gate two, the measured catch rate has to reach the ninety in a hundred we assumed in the table above, tested against your own labelled messages. Gate three, the Clock has to beat the baseline over a full in-season month. Gate four, nothing later is funded until the number it depends on exists.

A signed data agreement and consent from both of you, before any access. Read-only mailbox access through Google's own consent screen, revocable in one click, and no password ever shared. A forwarding rule works if you'd rather grant nothing, and here is

what it costs: it shows us what arrives but never when someone answered, so there is no median-minutes number and no second text when a guest has been waiting. Both need the read-only grant.

Your time, eight hours of it, priced into the table rather than called included. Fifteen minutes each on day one to agree the threshold and what counts as a cancellation. Reading the agreement. Signing off the reply wording. One call to your carrier. Ten minutes a week checking our labels. Then reading the verdict and deciding.

07

Four things that could make this the wrong call, and what we'd do about each

Every one of these is a real reason to say no.

Your existing vendor may do this for free. This is the biggest commercial risk here and we haven't retired it. FareHarbor Agent already covers email, and alerting an operator is an easier feature for them to add than the guest-facing one they've shipped. So we tell you to ask for the pilot now, before you spend anything with us.

The honest ceiling may be small. If two weeks turn up a handful of cancellation emails, the right answer is that you should build nothing. That exit is written in with a threshold agreed before any number exists, because a gate nobody would honour is worse than no gate. You keep the reply, you keep the count, and the money bought the number that stopped a larger mistake.

Late August only tells you about late August, so the finding gets labelled August-only rather than folded into a build to make the annual case look bigger.

We may be reading one sentence too hard. That FAQ line could be old policy text that no longer describes today. Two weeks of real messages tests it, and "the channel is fine" is a legitimate result that ends with us recommending you build nothing.

The next step is a reply, not a call

If this was useful, tell us. If it was wrong somewhere, tell us that instead, we'd rather know, and there are places above where we've guessed at your side of things. There is nothing to book and nothing to sit through. If you want the fortnight, the sooner it starts the more of your season it sees.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1 Your hours and the winter voicemail line
<https://www.northcountrycharters.com/>
- 2 The call-and-confirm request and your cancellation and weather rules
<https://www.northcountrycharters.com/faq>
- 3 Brian Ritchie and Kathy Brennan, and their published titles
<https://www.northcountrycharters.com/names-you-might-know>
- 4 One phone number and one shared inbox
<https://www.northcountrycharters.com/contact-1>
- 5 FareHarbor Agent and its select rollout
<https://fareharbor.com/blog/meet-fareharbor-agent-ai-powered-guest-communication-for-tour-operators/>
- 6 Emerald Air Service on text, WhatsApp and a 2027 portal
<https://www.emeraldairservice.com/>

- 7** Central Charters taking instant online booking
<https://fishingbooker.com/charters/view/41714>
- 8** Smokey Bay Air on FareHarbor
<https://smokeybayair.com/bear-viewing.html>
- 9** Destination Alaska Charters on FareHarbor
<https://destinationalaskacharters.com/>
- 10** Bear Viewing Alaska on FareHarbor
<https://www.bearviewingalaska.com/>
- 11** The 467-participant study on AI versus a human on bad news
<https://www.mdpi.com/2673-5768/6/1/36>
- 12** Your industry near the bottom of AI adoption
<https://www.fool.com/research/ai-adoption-statistics/>
- 13** AI review summaries downplaying complaints
<https://travel.yahoo.com/advice/safety/articles/booked-disaster-tripadvisor-ai-accused-223127463.html>