

FIELD STUDY PREPARED FOR TYONEK NATIVE CORPORATION

Your 22 federal registrations are an asset, and a missing signature page can take one of their records out of an evaluation.

ANCHORAGE, ALASKA AUGUST 2026

IN SHORT

We spent a day on Tyonek from the outside, using only public records, and this is what we would tell you if you had hired us.

The thing worth your attention is a single sentence in a solicitation you already know. Air Force RFP FA8773-25-R-0002 says affiliate past performance will be considered "provided that sufficient documentation is included in the proposal", then closes with the part that costs money: if that information is not included, the affiliate companies "shall not be evaluated or considered". The penalty is total, and the provision puts no duty on anyone to tell you it happened.

That sentence sits on an unusual structure. USASpending lists 22 child registrations under your parent UEI, with near duplicate names registered separately by state, including four whose names begin TFAB Ground Systems and three registered as Tyonek Engineering & Agile Manufacturing. Four of your companies carry SBA 8(a) certification. We can't see how your executed agreements are held from outside, which is part of what the half day checks.

Nothing we could verify shows Tyonek has ever failed to document an affiliate. GAO dismissed Tyonek Technical Innovations' protest in B-424364 on timing, not on merit, partly because the solicitation's language was, in GAO's words, unambiguously clear. That reads as close attention rather than as a slip. Our argument is about the size and the one way nature of the exposure, never about a failure we can't see.

So the recommendation is smaller than you would expect. Buy half a day, not a build. Our fee is \$4,500. With your people's time in it, the half day runs to about \$7,000. We model the decision itself as worth \$19,500 to \$42,900 against that, between \$2.75 and \$7.31 back for every dollar, and the table shows both ends, though part of that value is simply your not spending money with us. Let us measure four numbers off proposals you have already submitted. They are listed further down, and the one that matters most is how many of your bids even carry this provision.

The build, if the numbers ever justify it, is an indexed set of your executed agreements, a register of your 22 entities keyed on their registration numbers, and a rule that checks one against the other before you send. It never pays back in cash in any case we modelled, and it only approaches it at 22 qualifying bids a year, a number we chose and not your registration count.

We have written a kill condition into this study, before any money moves. If all three of these hold, completeness already above about 90 percent, assembly under a day a bid, and fewer than about eight qualifying bids a year, we will tell you in writing not to build it.

01

One sentence in an Air Force solicitation you already bid decides whether an affiliate's record counts

The provision is short, we have verified it in one of your solicitations, and it carries no duty to tell you when it bites.

Air Force RFP FA8773-25-R-0002, issued 19 May 2025 as an 8(a) set aside, spells out how affiliate past performance gets treated. Affiliate, sister, teaming and joint venture arrangements are considered when sufficient documentation rides along in the proposal, and the prime must demonstrate the affiliate will perform significant and critical aspects of the contract. Then the sentence that carries the risk: "If the required information is not included in the proposal, the proposed affiliate companies, sister companies, joint venture companies, etc. shall not be evaluated or considered."

Read that as an operations problem rather than a legal one. An affiliate whose paperwork is short does not get marked down. It stops existing for that evaluation. The work was done, the rating is on file, and none of it counts.

22

child registrations under your parent UEI, each one a potential source of past performance and a potential documentation gap

What makes this yours specifically is the shape of the corporation underneath that sentence. USASpending lists 22 child recipient registrations under parent UEI FM96SF3VF8H9, and the names are close enough to trip a person in a hurry. Four registrations begin TFAB Ground Systems, spread across Alaska, Alabama twice and Georgia. Three are registered as Tyonek Engineering & Agile Manufacturing across Georgia, Alabama and Alaska. Four subsidiaries hold SBA 8(a) certification, including Tyonek Technical Services, which holds the \$98,115,620 Cyber Operations Formal Training Support award.

We should say plainly what we did not find. There is no public evidence that Tyonek has ever submitted an affiliate without its documentation. GAO's dismissal in B-424364 rested on timing and prematurity, and it recorded that Tyonek remained in the competition. This is a control worth having because the provision carries no obligation to warn you.

A struck affiliate is a total loss, and nothing obliges anyone to tell you it happened

Most process problems cost you a percentage. This one costs you the whole item.

When an affiliate is struck for missing documentation, the past performance you meant to lean on is gone from that evaluation and the evaluator sees a thinner company than the one that exists.

The provision imposes no duty on the agency to tell you. Whether it can be cured depends on whether the agency opens discussions, which you can't count on.

25

locations in 16 states, with manufacturing at Madison, Huntsville and Warner Robins and repair stations at Kiln and Huntsville

We are not going to tell you what this costs you per year, because we can't know it and the number would be invented. The cost is the value of an affiliate's record on a bid where it mattered, multiplied by however often that has actually happened, and that second number is the one we have no way to see from outside, which is why one of the four numbers the half day measures is whether any debrief has ever recorded an affiliate that was not evaluated.

03

What changes is that the check happens before you send

The work does not disappear. It moves.

Here is one real task, assembling the affiliate documentation for a single bid. We have not seen how you do this, so the Today column is the sequence the provision forces on

anyone, not a description of your process. Replacing it with your actual steps is the first thing the half day does.

Assembling affiliate past performance for one proposal

TODAY	AFTER
1 Decide which affiliates' past performance to cite	Decide which affiliates' past performance to cite. Unchanged, this is a capture judgement.
2 Establish which legal entity holds each record, across 22 registrations with near duplicate names	The entity is resolved by registration number, on UEI and CAGE. Names are never matched, so two TFAB entities can't be confused.
3 Hunt for the executed affiliation, teaming or JV agreement wherever it is held	The agreement is retrieved from one indexed set, with the signature page on screen.
4 Read the provision and confirm what it demands	Unchanged, we assume a capture manager does this in about ten minutes
5 Check by hand whether each required document is actually in the proposal	A rule reports present, missing, expired or ambiguous per entity and per document. Same answer every time.
6 Nothing checks the assembled set as a whole before it goes out	A named person clears a one page gap list and signs for it before submission.
7 Judge and state that the affiliate will perform significant and critical aspects of the contract	Judge and state that the affiliate will perform significant and critical aspects of the contract. Deliberately unchanged.

Two things to say about this figure. The Today column is assumed, not observed, and the half day tests it. The last row is deliberately identical, because that judgement is a contractual representation about future work share and it belongs to a person at Tyonek. A system that made that call for you would be a liability, not a feature.

Most of what we would build is a filing system and a rule, and the rule is where we refuse to put a model


Most of the value here is a filing system and a rule.

Four pieces. A collected set of your executed affiliation, teaming and joint venture agreements, gathered from wherever they live now and indexed. A register of your 22 entities keyed on their registration numbers. A rule that compares what a solicitation demands against what is actually on file. Finally, a one page gap list a named person signs before the proposal goes out.

On a given bid, the rule takes the affiliates you intend to cite and reports, for each one and each required document, whether it is present, missing, expired or ambiguous. Anything not verified against a page in a document reports as missing.

Two places use a model, and neither can change that verdict. The first reads a scanned agreement as it is filed and proposes what it is, who signed it and when it expires, refusing to answer without pointing at the page. A person confirms that once, at filing, and the confirmed fact is what the rule reads later. The second drafts the affiliate narrative from facts a person has already approved.

A third use is possible, a second reading of the solicitation to flag where it disagrees with yours. We left it out, because reading the provision takes a capture manager about ten minutes, which we are assuming rather than measuring, and a second read does not earn its keep at your bid volume.



Read left to right. Two things meet at the check in the middle, your filed agreements along the top and the bid you are about to send. The check is a rule, not a model. The reading step is confirmed by a person before anything it produced can be used, and the drafting step only runs after a person has signed. Neither model can move a go or no go.

The check is a list compared against a list, not a prediction. It gives the same answer every time and can be tested exhaustively, which is why we would not put a model near it. Entity matching runs on registration numbers only, never on a name.

There are three places we would refuse to go. Anything touching a maintenance sign off or return to service, because 14 CFR 145.213(d) reserves that to a certificated mechanic or repairman. Anything drawing-derived, including AS9102 first article content, if it is export controlled, because a hosted model has to decrypt your data to read it and we read 22 CFR 120.54 as putting that outside the carve out. Your government property records are the third, because they are rule governed and structured, and you posted an analyst in July to own them.

We would also not recommend predictive maintenance to you. GAO reports the services "generally lack metrics to evaluate the results of predictive maintenance". The DoD Inspector General found the department had not fully implemented it on any existing weapon system. Your capability page mentions no AI at all, and on this particular thing that is the right answer, not a gap.

Buy the storage, the identity, the text recognition for scanned pages, and the model itself. Never train one. What is worth building is narrow: the register keyed on your registration numbers, the rule set and the cases that prove it, and a check inside that first reading step that keeps export marked documents away from any outside service.

If you already run a system holding these agreements, extending it beats standing up anything new.

05

The half day pays for itself. The build does not pay back in cash in any case we can model

Two things are priced here, and only one of them is the ask.

Two things are priced. The half day we are proposing, whose product is a decision. Second, the build that half day would decide about, which we priced but are not asking you to buy. Every figure is computed from the drivers in the footer except where a line is marked assumed. Three are, the bid volume, what the decision is worth, and the return per dollar that falls out of the two.

Both layers, three scenarios each. The marks show where each number comes from.

	Conservative	Most likely	Aggressive
 Qualifying bids a year, assumed	8	14	22
 The half day, total cost to you	\$7,080	\$6,095	\$5,865
 What the decision is worth to you	\$19,500	\$31,200	\$42,900
 What comes back for every dollar	\$2.75	\$5.12	\$7.31
 The build, five year cost	\$213,600	\$171,100	\$140,300
 What the build returns in cash over five years	\$7,650	\$30,960	\$109,620
 Share of the build's cost recovered in cash	4%	18%	78%
 Hours the build frees over five years	158	674	2,201
 Months to pay back the build in cash	never	never	never

The half day is our fee of \$4,500 plus your own people's time, \$1,400, \$800 and \$600 across the three columns, plus contingency at 20, 15 and 15 percent. That is why the conservative column costs more than the aggressive one for the same half day. \$4,500 is the only invoice. Drivers, so you can check the arithmetic. Capture staff at \$85 an hour fully loaded and contracts administration at \$60, both assumed by us because we have no view of your payroll, and both held constant across scenarios so there is one number to argue with rather than three. Qualifying bids of 8, 14 and 22 a year. Those three are ours, not yours, and they are not derived from anything, which is why the row is marked assumed, and why counting them is one of the four numbers the half day measures and the one the whole case scales on. The 22 is not your registration count, the resemblance is a coincidence. Assembly of 12, 20 and 32 hours a bid, cut by 20, 35 and 50 percent, which is a day and a half at the low end and four days at the high end. Finding a named agreement, 5, 7 and 10 hours a bid, cut by 45, 60 and 70 percent. A first year ramp of 25, 30 and 35 percent for a mid year start, then full run rate, with no compounding modelled. Contingency of 20, 18 and 15 percent on the build. The build's cost is implementation plus five years of running it. In the most likely case that is about \$78,000 to implement, \$8,500 to train, and \$13,000 a year to run for four and a half years, plus contingency. The conservative column is higher than the aggressive one on every one of those lines, not on contingency alone, because we pad what we are least sure of. The five year cost of \$213,600 is about \$43,000 a year, call it between a quarter and a third of a person. The risk line, in full, because the endpoints alone are not enough to check it. Events a year 0.3, 0.6 and 1.2, which is about once every three years at the low end and a little more than once a year at the high end. Value per event \$12,000, \$20,000 and \$30,000, being proposal work that stopped counting, never a contract value. Share of those events the check would actually catch before submission, 50, 60 and 70 percent. That last factor is the one that turns the other two into the cash row, and it is our estimate, because a completeness check catches gaps in documents and not every way a bid can go wrong. Multiply the three, apply the ramp, and the conservative case is \$1,800 a year and \$7,650 over five. Cash and hours are kept apart on purpose. The 2025 Top 49ers listing records Tyonek at 1,523 worldwide and 15 statewide, and we have no way to know who touches a federal proposal or where they sit, so we will not convert freed hours into money. The hours row is hours. 2,201 hours over five years is roughly one person year in total, about a quarter of someone's year in each of the four full years. 158 hours is about four weeks of one person across the whole five. What the decision is worth is our assumption and not a computation from your data, which is why it is marked that way. We put it at the chance the measurement changes your decision, roughly 1 in 4, 2 in 5 and just over half, multiplied by \$78,000. That \$78,000 is the Phase 1 implementation in the most likely case, which is what you would commit first if you went ahead, and it is smaller than the five year cost above because a decision to stop at the gate never incurs the run cost. It is partly circular by construction, because part of what the half day buys is not spending money with us. The part that is not circular is that you keep the four numbers whether or not you build anything, and whether or not you hire us again. The conservative column assumes exactly

8 qualifying bids, which sits on one of our three kill conditions. That is deliberate. It is the lowest bid volume at which building this is arguable at all, holding the other two conditions open.

 modelled from stated drivers  assumed, and named as such

never

is when the build pays back in cash, in every case we can model. The half day pays back in months.

Most projects like this do not work. MIT's NANDA review found roughly 95 percent of enterprise AI pilots produce no measurable financial return, and RAND puts outright AI project failure at about 80 percent. Both are directional rather than precise, the denominators are debated, and the MIT figure measures pilots without profit and loss impact rather than software that does not run. Oliver Wyman surveyed more than 150 aviation professionals and found 58 percent reporting their companies stuck at the experimental stage, unchanged year over year. Three things make this different, and none of them is enthusiasm. There is very little model in it, so there is little to misbehave. The measurement comes before the money. The one mistake that would cost you a bid, matching two of your own companies by name, is impossible by construction here, because the register only ever reads registration numbers.

The owner should be whoever runs the pursuits that cite affiliate past performance, and the same four numbers get measured again after one bid cycle, against the baseline, by that person. If no single person owns them after the half day, that is a reason not to build.

We would measure four numbers first, and everything after them is conditional on what they say

The half day is a review of your last five to ten completed bids that named affiliate past performance, using documents you already hold. Nothing is installed and nothing leaves your building. It produces four numbers: how many affiliates shipped with a complete documentation set, how many hours assembly actually took, how many bids in the last year even qualify, and whether any debrief has ever noted an affiliate that was not evaluated.

The third matters most, because the whole case scales on it. We have verified this provision in one solicitation, and how many of your bids carry something like it is the number we are least able to guess.

Phase	What runs	How we would know it worked
Now	The half day baseline review, on proposals you have already submitted	Four measured numbers, and a written verdict within five working days
	Confirm someone owns the completeness number and is in the room for the readout	One named person, or we stop here
	Spot check how many signed agreements can actually be put hands on in an afternoon	Share of required documents located, and how many places they came from
Next	Collect and index the executed agreements wherever they are held. This is the real work and it is not AI work	Share of required documents indexed, and time to retrieve a named agreement
	Build the entity register across all 22 registrations, keyed on registration numbers only	22 of 22 resolved by identifier, zero name matches anywhere
	Run the completeness rule alongside live bids before it becomes a required step	Rule verdicts compared against your capture manager's own read
	Set the bar for the page reader before it is trusted to prefill anything	Its proposed reading accepted unchanged by the person confirming on at least 9 of 10 documents, or it stays a suggestion only
	Hand over the rule set, the register and the runbook	One full bid cycle run by your team without us

Five gates, and each one is your exit. Before anything, a named owner exists. After the half day, the kill fork: if all three of these hold, completeness already above about 90 percent, assembly under a day a bid, and fewer than about eight qualifying bids a year, we write down that this should not be built and we do not build it. Then a do-the-documents-exist gate, because if the signed agreements can't be found, the honest answer is a records fix rather than software. Then rules before models, with the rule proved against every case we can write down before any model work is funded. Then the check runs alongside live bids without being required, before it becomes required.

Phase	What runs	How we would know it worked
	Re-measure the same four numbers after one full bid cycle	Completeness rate and assembly hours against the half day baseline, in writing
Later	Registration currency and 8(a) status watching across the 22 entities. A lapse costs you eligibility with no warning attached, runs on public data, and is honestly not an AI build	Lapse conditions caught before they cost anything
	Government property records, once that role is filled and whoever holds it has defined the record set	Exceptions caught before an audit rather than during one
	Quality records and first article content, which reopens only behind a model running inside your own environment	Not costed until that condition is met

Five gates, and each one is your exit. Before anything, a named owner exists. After the half day, the kill fork: if all three of these hold, completeness already above about 90 percent, assembly under a day a bid, and fewer than about eight qualifying bids a year, we write down that this should not be built and we do not build it. Then a do-the-documents-exist gate, because if the signed agreements can't be found, the honest answer is a records fix rather than software. Then rules before models, with the rule proved against every case we can write down before any model work is funded. Then the check runs alongside live bids without being required, before it becomes required.

Half a day of whoever assembled your last few bids, the submitted proposals themselves, and confirmation of who owns federal capture. That is the entire ask. Later, if the numbers justify it, read access to the executed agreements, a ruling from contracts on where those documents may live, and confirmation of which entities belong in the register.

Four things could make us wrong, and the first is that you may not have this problem

Four things, and the first is the one that would cost you the most to discover late.

The first is that this may not be a problem you have. Your completeness rate might already be at or near 100 percent, in which case the build is worth far less than we modelled and possibly nothing. That is exactly why the ask is half a day of measurement with a kill condition attached.

The second is that our confidence is genuinely moderate. We put the chance that this is a live gap at your company at a bit under a coin flip, 0.4 to 0.5, and the reason cuts against us: GAO dismissed Tyonek Technical Innovations' protest partly because the provision's language was unambiguously clear, which is a point in favour of a capture team complying with it unaided. That is a different question from the one priced further down, where we ask how often the measurement actually changes what you decide and put that at 1 in 4 to just over half.

The third is that the hardest part is not the clever part. Finding executed agreements across a company operating in 16 states can take far longer than anyone estimates, and some may not be findable at all. We priced that work inside the scope rather than calling it your homework.

The fourth is ownership. Your headquarters is in Anchorage while your manufacturing sits in Madison, Huntsville and Warner Robins and your repair stations in Kiln and Huntsville, so the people who would use this are spread across several states from the people who would fund it. Your about page names Joe Webb as Vice President, Business Development, and we have no way to see from outside whether federal capture sits with him. A build with nobody clearly owning it is the most reliable way for this kind of project to quietly stop, so the first conversation should settle who owns the completeness number. If nobody does, we would say so and recommend against starting.

The next step is small

Half a day with whoever assembled your last few bids, using proposals you have already sent. You keep the four numbers and the written verdict whichever way they fall, including the version where we tell you not to build this.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1 GAO decision B-424364, the affiliate past performance provision quoted in full, the 19 May 2025 issue date, the 8(a) set aside under FAR subpart 15.3, and the dismissal on timing and prematurity
<https://www.gao.gov/products/b-424364,b-424364.2>
- 2 22 child registrations under parent UEI FM96SF3VF8H9, including the near duplicate entity names
<https://api.usaspending.gov/api/v2/recipient/children/FM96SF3VF8H9/>
- 3 Tyonek leadership including Joe Webb as Vice President of Business Development, 25 locations in 16 states, and the two FAA Part 145 repair stations
<https://www.tyonek.com/about/>
- 4 Certifications, customer set, manufacturing footprint and employee statement
<https://www.tyonek.com/capabilities/>
- 5 The four SBA certified 8(a) companies
<https://www.tyonek.com/service-companies/>
- 6 The \$98,115,620 Air Force Cyber Operations Formal Training Support award
<https://madisonidb.com/tyonek-technical-services-wins-98-1m-u-s-air-force-contract/>
- 7 The Coast Guard voice and flight data recorder award
<https://www.tyonek.com/2024/10/10/tyonek-aircraft-overhaul-llc-secures-contract->

[with-u-s-coast-guard-for-voice-and-flight-data-recorder-replacement/](#)

- 8 2024 revenue and the worldwide and statewide employee counts
<https://digital.akbizmag.com/issue/october-2025/2025-top-49ers/>
- 9 Quality inspector duties including AS9102 and non conformance reports
<https://tyonekjobs.applicantpro.com/iframe/4144030.html>
- 10 The Government Property Analyst posting, dated 14 July 2026, and its FAR Part 45 requirement
<https://tyonekjobs.applicantpro.com/iframe/4146846.html>
- 11 14 CFR 145.213(d), who may sign off a maintenance release
<https://www.law.cornell.edu/cfr/text/14/145.213>
- 12 14 CFR 43.9, the signature as approval for return to service
<https://www.law.cornell.edu/cfr/text/14/43.9>
- 13 22 CFR 120.54, the ITAR end to end encryption carve out
<https://www.law.cornell.edu/cfr/text/22/120.54>
- 14 GAO on the military services lacking metrics for predictive maintenance
<https://www.gao.gov/products/gao-23-105556>
- 15 DoD Inspector General on predictive maintenance not fully implemented
<https://fedscoop.com/watchdog-dod-still-hasnt-fully-deployed-predictive-maintenance-on-any-weapon-systems/>
- 16 Oliver Wyman, 58 percent of aviation professionals report being stuck at the experimental stage
<https://www.oliverwyman.com/our-expertise/insights/2026/apr/aviation-mro-labor-and-material-supply-chain-paradigm.html>
- 17 MIT NANDA, the share of enterprise AI pilots showing no measurable financial return
https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf
- 18 RAND on the rate at which AI projects fail
https://www.rand.org/pubs/research_reports/RRA2680-1.html
- 19 The Stennis overhaul facility's ISO 9001:2015, AS9100 Rev D, AS9110 Rev C and FAA Part 145 certification
<https://www.tyonek.com/reusable-block/services-tyonek-facilities/>

